

# **Online matching with graph neural networks**

**Ellen Vitercik** (Stanford)

# Machine learning for optimization

When designing algorithmic solutions to computational problems in practice:

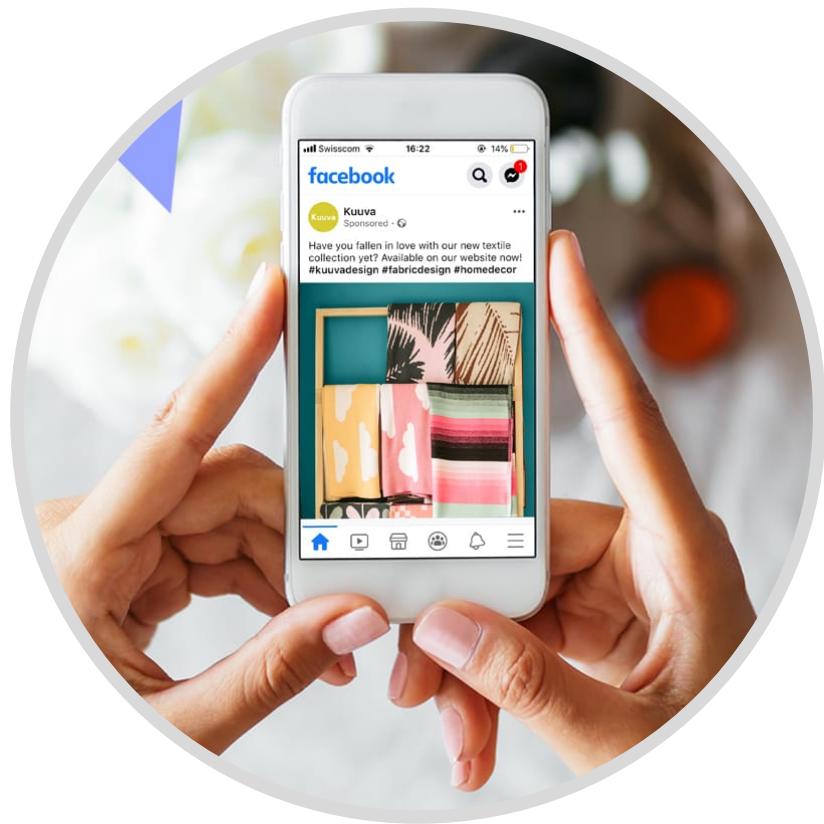
- Often **ample data** about application domains



Routing



Rideshare



Digital advertising



Scheduling

- Data can potentially be harnessed to **optimize algorithmic performance**

# Why machine learning?

Problems that a shipping company solves to route its trucks will change daily  
...but not drastically

Demand and traffic will vary  
...but road network will remain the same



Likely **underlying structure** that can be uncovered with the help of ML

**Goal:** use structure to optimize algorithm performance on future problems

---

Runtime, solution quality, ...

# Rich history

## **Data-driven approaches to algorithm selection**

[e.g., Lobjois, Lemaître, '98; Gomes, Selman, '01; Leyton-Brown et al., 2009, Kadioglu et al., 2010, Sandholm, 2013, ...]

## **ML for algorithm configuration**

[e.g., Hutter et al. '09, '10, '11, Ansótegui et al. '09, Sandholm, 2013]

## **Integration of modern ML models into discrete solvers**

[e.g., surveys by Bengio et al., '18; Cappart et al., '23; ...]

## **Data-driven robust optimization**

[e.g., Esfahani, Kuhn, '17; Bertsimas et al., '18; ...]

## **Theoretical guarantees**

[e.g., book chapters by Balcan, '20; Mitzenmacher, Vassilvitskii, '20; ...]

# Key challenges

ML approaches to discrete optimization **rarely work out of the box**:  
must be **aligned** to the algorithmic task at hand

## Major challenges:

- **Architecture:** seq2seq, GNNs, transformers, hyperparameters, ...?
- **Supervision:** how to supervise on (NP-hard) algorithmic problems?
- **Robustness:** can we provide any guarantees?

# Outline

1. Introduction

**2. Graph neural networks (GNNs) for online matching [ICML'24]**



Alexandre Hayderi



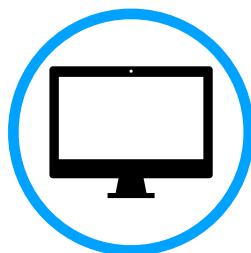
Amin Saberi



Anders Wikum

# Online matching

## Critical problem in many digital marketplaces



**Advertising:** website visitors matched to ads

[e.g., Mehta et al., JACM'07]



**Crowdsourcing platforms:** crowdworkers matched to tasks

[e.g., Tong et al., VLDB'20]



**Rideshare:** riders matched to drivers

[e.g., Zhao et al., AAAI'19]



**Medicine:** organ donors matched to patients

[e.g., Ezra et al., EC'20]

**Primary challenge: irrevocable** matching decisions must be made online  
**without** precise knowledge of how demand will evolve

# Our contributions

## GNNs for **Online Bayesian Bipartite Matching (OBBM)**

**Theoretical:** justify the suitability of GNNs for OBBM

**Empirical:** train a GNN for **Online Bayesian Bipartite Matching (OBBM)**

- Supervise with actions of **optimal online algorithm**  $\text{OPT}_{\text{on}}$
- Outperforms baselines on synthetic/semi-synthetic graphs

**Key challenge:** sheer complexity of  $\text{OPT}_{\text{on}}$ !

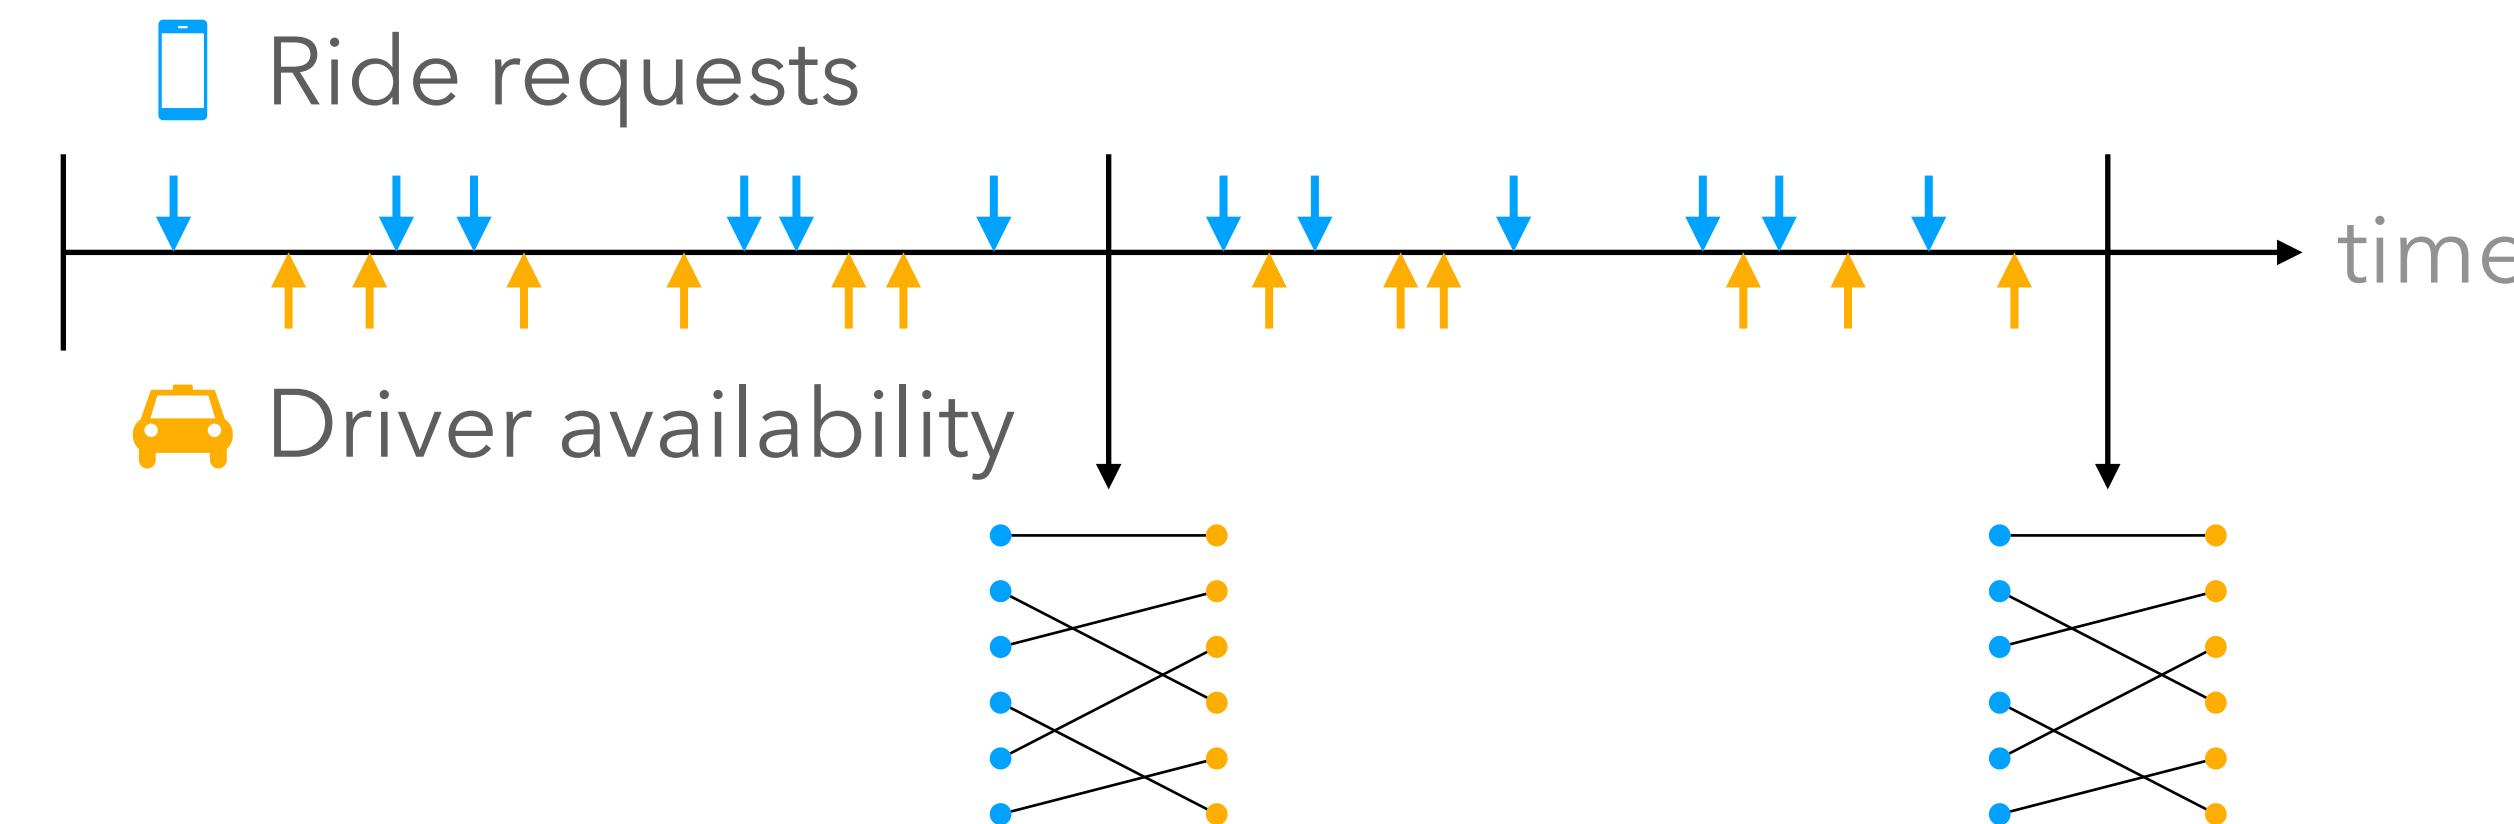
NP-hard to provide an  $\alpha$ -factor approximation to  $\text{OPT}_{\text{on}}$  (for some constant  $\alpha$ )

[Papadimitriou et al., EC'21]

# Related work: RL for rideshare

Xu et al., KDD'18; Qin et al., INFORMS J. Appl. Anal. '20; Azagirre et al., INFORMS J. Appl. Anal. '24; ...

- Open trip requests and available drivers are batched within time windows
- Matched at predefined discrete review times



- In contrast, we aim to match as soon as request arrives  
*E.g., Adwords [Mehta et al., JACM'07]*
- Future work: possible to combine techniques?

# Related work: ML for online matching

Focus on purely **theoretical** [e.g., Antoniadis et al., ICML'20, NeurIPS'20]

**Adversarial learning** to find good worst-case alg [e.g., Kong et al., ICLR'18]

**Most related:** Almorani et al. [TMLR'22]

- **RL** approach for unknown i.i.d. setting (optimal online algorithm unknown)
- In OBBM, can use supervision from  $\text{OPT}_{\text{on}}$  and **train on much smaller graphs**
- Learns a **tailored matching policy** per graph configuration, unlike us

Li et al. [ICML'23] bridge the gap between **worst-case/average-case**:

- **Dynamically switch** between expert and ML predictions
- Future work: combine with our approach for **robustness** guarantees?

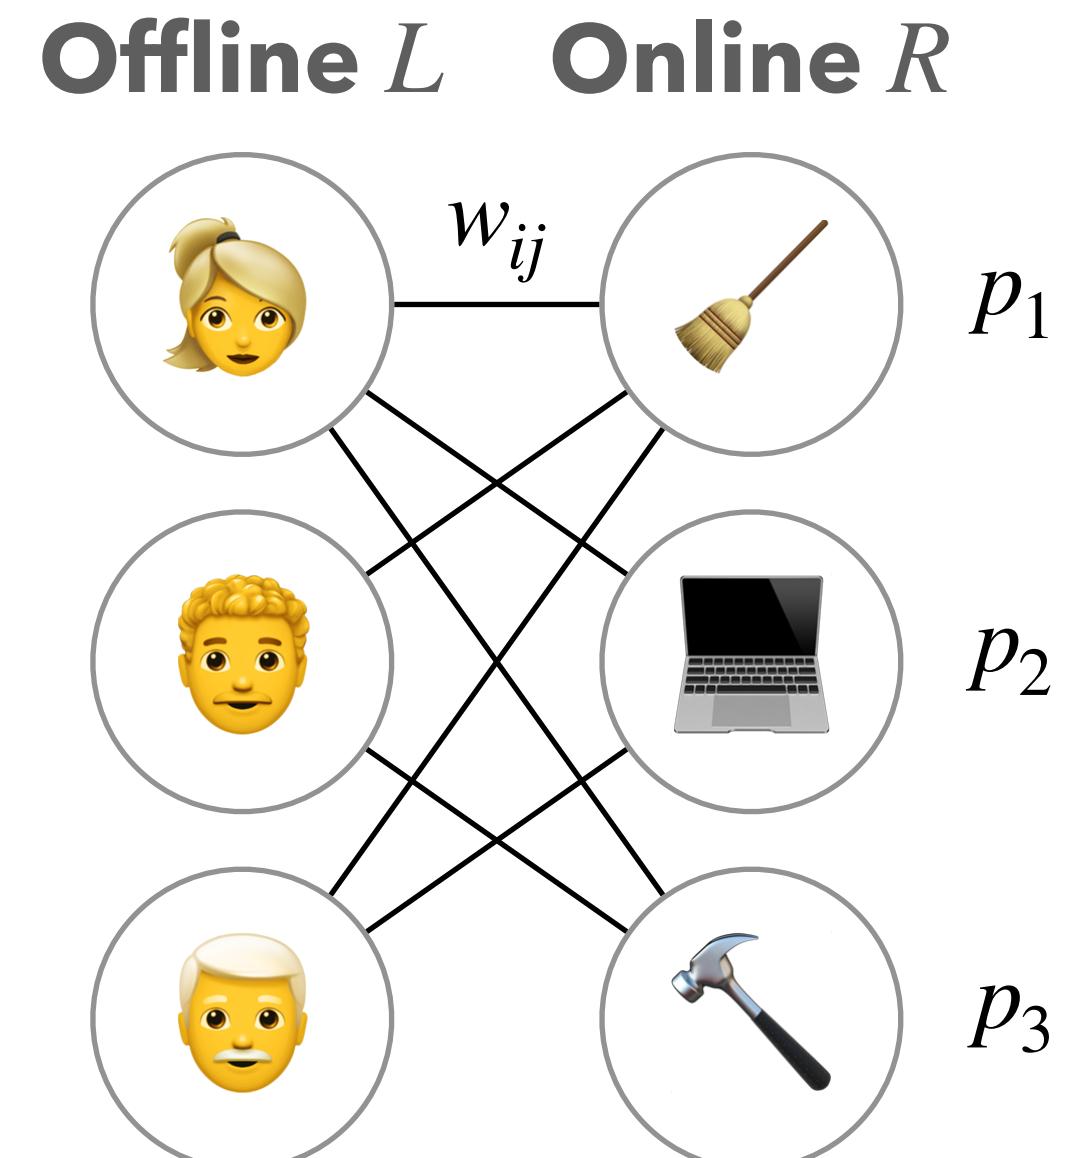
# Outline

1. Introduction
2. Graph neural networks (GNNs) for online matching [ICML'24]
  - i. **Background**
    - a. **Online Bayesian bipartite matching**
    - b. GNNs
  - ii. GNN setup
  - iii. Experiments
  - iv. Theoretical guarantees

# Online Bayesian Bipartite Matching (OBBM)

E.g., Krengel, Sucheston, Probab. Banach Spaces '78; Papadimitriou et al., EC'21; Braverman et al., EC'22;...

- Bipartite graph  $G = (L, R, E)$  with offline nodes  $L$ , online nodes  $R$ 
    - E.g., represent **crowdworkers** and **tasks**
  - Edge weights  $w_{ij}$  for each  $(i, j) \in E \subseteq L \times R$ 
    - E.g., represent **workers' payoffs** for completing tasks
  - Online node arrival probability  $p_t \in [0,1]$  for each  $t \in R$
  - If node  $t$  appears, irrevocably decide:
    - Match  $t$  with an unmatched offline neighbor, or
    - Skip  $t$  and not match it to any node
- Goal:** compute a high-weight matching



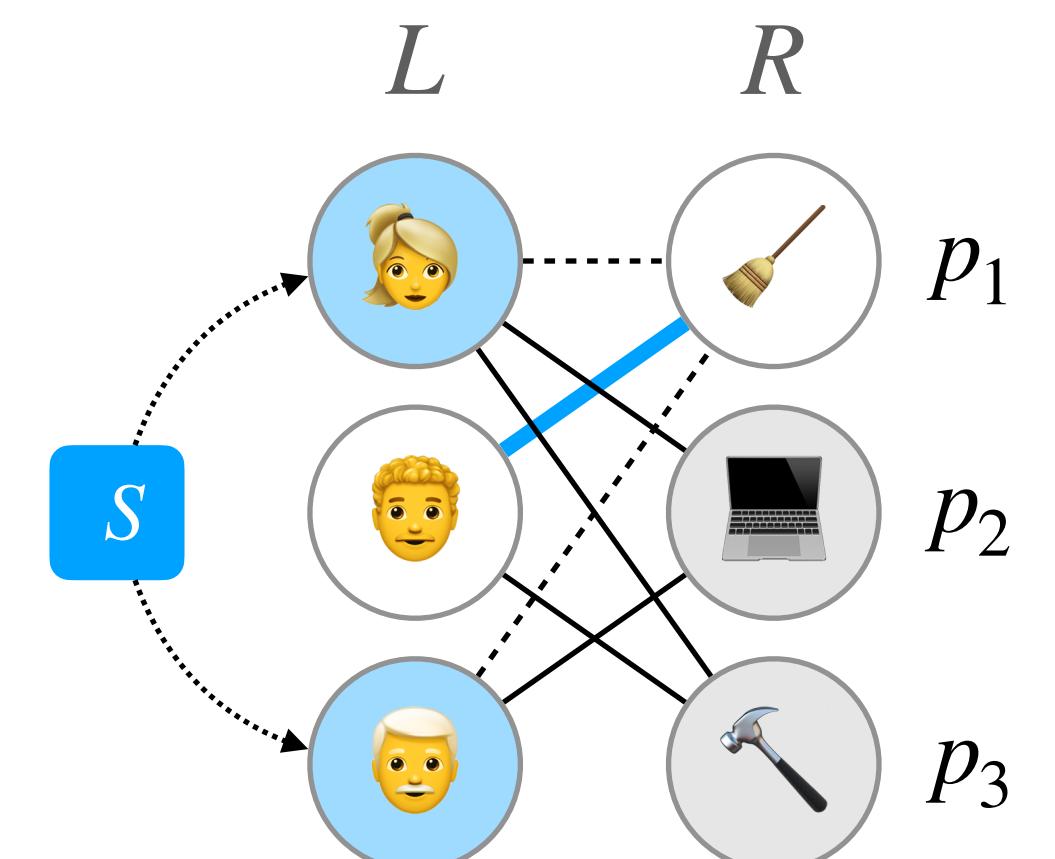
# Online optimal algorithm and value-to-go

Let  $S$  be the set of offline nodes available on round  $t$

$\text{OPT}_{\text{on}}$  computes the **value-to-go (VTG)** function (or *Bellman equation*)  $V_G(S, t)$

$V_G(S, t)$  = expected value of the maximum weight matching achievable on  $G$

- Expectation over arrivals  $\{t, \dots, |R|\}$
- Matchings restricted to  $S$



# Online optimal algorithm and value-to-go

Let  $S$  be the set of offline nodes available on round  $t$

$\text{OPT}_{\text{on}}$  computes the **value-to-go (VTG)** function (or *Bellman equation*)  $V_G(S, t)$

$V_G(S, t)$  = expected value of the maximum weight matching achievable on  $G$

- Expectation over arrivals  $\{t, \dots, |R|\}$
- Matchings restricted to  $S$

$$V_G(S, t) = \underbrace{(1 - p_t)}_{\downarrow} \cdot V_G(S, t + 1) + p_t \cdot \max \left\{ V_G(S, t + 1), \max_{u \in N_G(t) \cap S} \left\{ w_{tu} + V_G(S \setminus \{u\}, t + 1) \right\} \right\}$$

If  $t$  doesn't arrive, don't get any value from node  $t$

# Online optimal algorithm and value-to-go

Let  $S$  be the set of offline nodes available on round  $t$

$\text{OPT}_{\text{on}}$  computes the **value-to-go (VTG)** function (or *Bellman equation*)  $V_G(S, t)$

$V_G(S, t)$  = expected value of the maximum weight matching achievable on  $G$

- Expectation over arrivals  $\{t, \dots, |R|\}$
- Matchings restricted to  $S$

$$V_G(S, t) = (1 - p_t) \cdot V_G(S, t + 1) + p_t \cdot \max \left\{ \underbrace{V_G(S, t + 1)}_{\downarrow}, \max_{u \in N_G(t) \cap S} \left\{ w_{tu} + V_G(S \setminus \{u\}, t + 1) \right\} \right\}$$

If  $t$  does arrive      We can skip it, and get no value from node  $t$

# Online optimal algorithm and value-to-go

Let  $S$  be the set of offline nodes available on round  $t$

$\text{OPT}_{\text{on}}$  computes the **value-to-go (VTG)** function (or *Bellman equation*)  $V_G(S, t)$

$V_G(S, t)$  = expected value of the maximum weight matching achievable on  $G$

- Expectation over arrivals  $\{t, \dots, |R|\}$
- Matchings restricted to  $S$

$$V_G(S, t) = (1 - p_t) \cdot V_G(S, t + 1) + p_t \cdot \max \left\{ V_G(S, t + 1), \max_{u \in N_G(t) \cap S} \left\{ \frac{w_{tu} + V_G(S \setminus \{u\}, t + 1)}{\downarrow} \right\} \right\}$$

Or match to some unmatched neighbor  $u \in N_G(t) \cap S$ , get value  $w_{tu}$ , and remove  $u$  from unmatched set  $S$

# Online optimal algorithm and value-to-go

Let  $S$  be the set of offline nodes available on round  $t$

$\text{OPT}_{\text{on}}$  computes the **value-to-go (VTG)** function (or *Bellman equation*)  $V_G(S, t)$

$V_G(S, t)$  = expected value of the maximum weight matching achievable on  $G$

- Expectation over arrivals  $\{t, \dots, |R|\}$
- Matchings restricted to  $S$

$$V_G(S, t) = (1 - p_t) \cdot V_G(S, t + 1) + p_t \cdot \max \left\{ V_G(S, t + 1), \max_{u \in N_G(t) \cap S} \left\{ w_{tu} + V_G(S \setminus \{u\}, t + 1) \right\} \right\}$$

If node  $t$  arrives, don't match if  $\square > \square$ , else match to maximizing  $u$

# Outline

1. Introduction
2. Graph neural networks (GNNs) for online matching [ICML'24]
  - i. Background
    - a. Online Bayesian bipartite matching
    - b. GNNs**
  - ii. GNN setup
  - iii. Experiments
  - iv. Theoretical guarantees

# GNN architecture

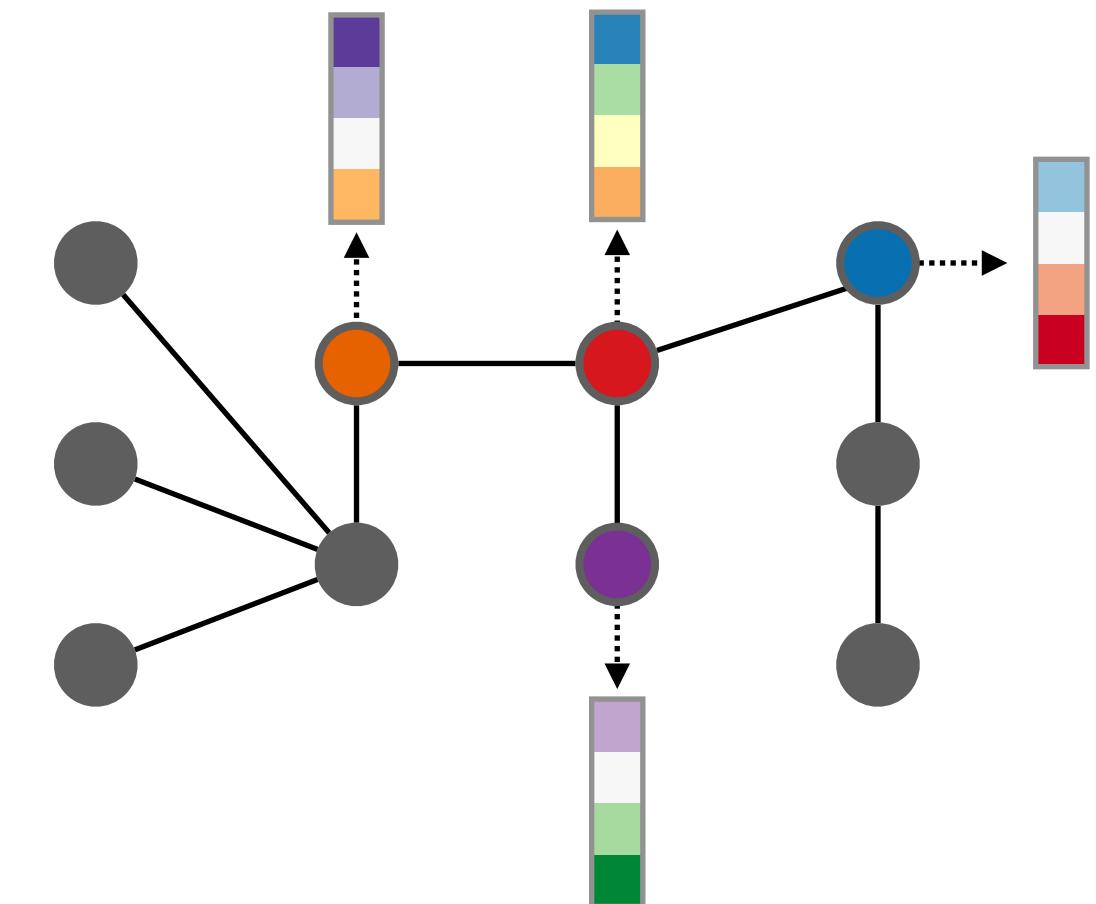
**GENConv** [Li et al. '20]

Each node  $v$  begins with initial embedding  $\vec{h}_v^{(0)}$  (input features)

On iteration (or “layer”)  $k \in \{1, \dots, L\}$ , for each node  $v$ :

1. Aggregate neighbors’ embeddings  $\vec{m}_v^{(k)} = \max_{u \in N(v)} \left\{ \text{MLP} \left( \vec{h}_u^{(k-1)}, w_{vu} \right) \right\}$
2. Update node embedding  $\vec{h}_v^{(k)} = \text{MLP} \left( \vec{h}_v^{(k-1)} + \vec{m}_v^{(k)} \right)$

Use  $\vec{h}_v^{(L)}$  to make predictions about node  $v$

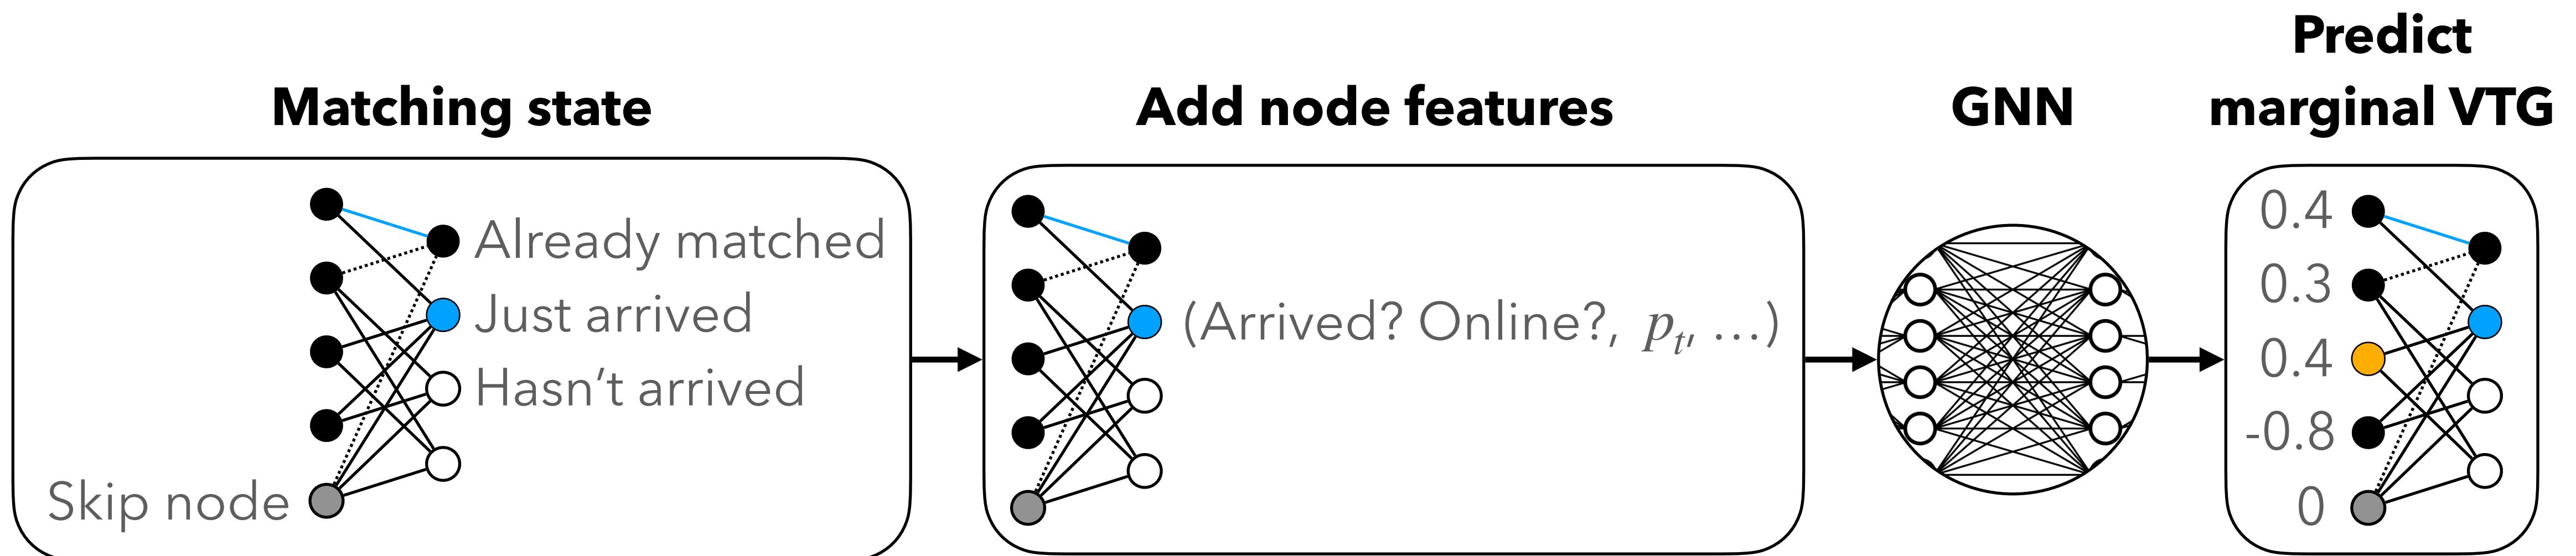


# Outline

1. Introduction
2. Graph neural networks (GNNs) for online matching [ICML'24]
  - i. Background
  - ii. GNN setup**
  - iii. Theoretical guarantees
  - iv. Experiments

# MAGNOLIA

Matching Algorithms via **GNNs** for **Online** Value-to-go **A**pproximation



$$\text{Supervise with marginal VTG} \left( w_{tu} + V_G(S \setminus \{u\}, t+1) \right) - V_G(S, t+1)$$

Value of matching with node  $u$       Value of skipping

# Outline

1. Introduction
2. Graph neural networks (GNNs) for online matching [ICML'24]
  - i. Background
  - ii. GNN setup
  - iii. Theoretical guarantees**
  - iv. Experiments

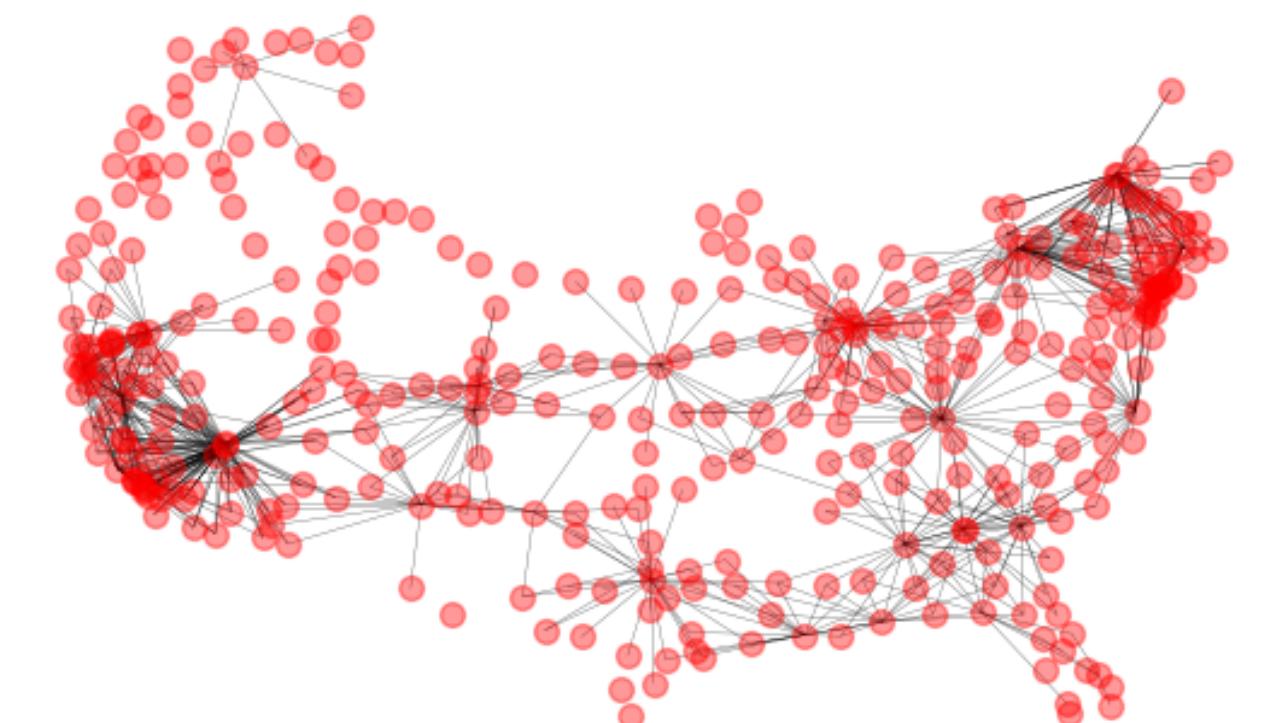
# Why are GNNs well-suited for online matching?

Most practical instances are “well-structured”

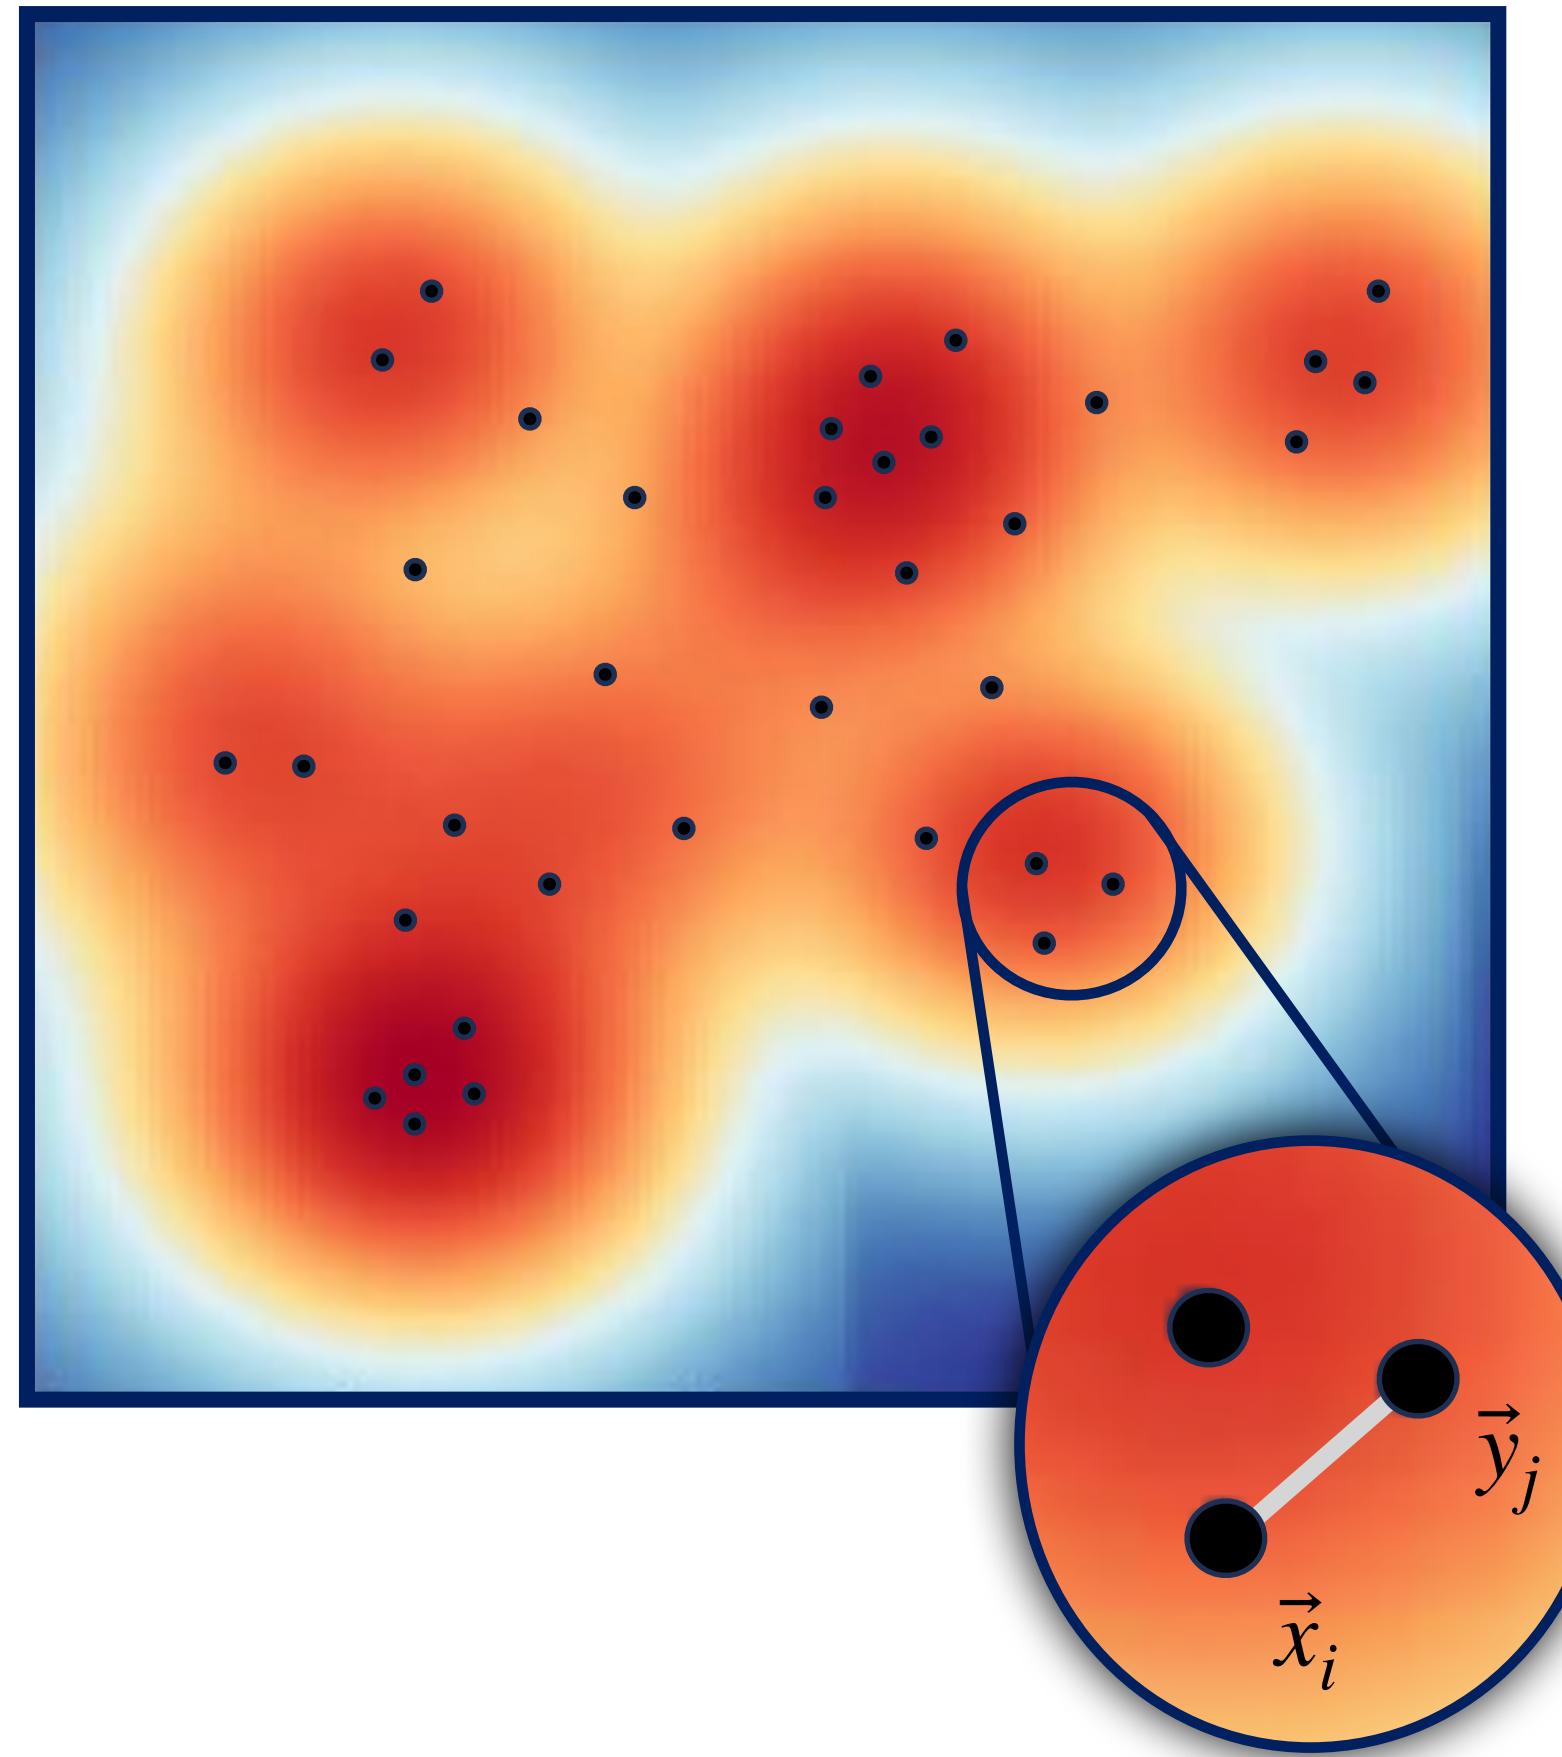
We show this rigorously for **random geometric graphs**

*We prove that under these families, **VTG is nearly local***

Already known that **GNNs compute local functions**



# Random geometric graphs $\mathcal{G}(|L|, |R|, D, \Delta)$



Draw  $\vec{x}_1, \dots, \vec{x}_{|L|} \sim D$  and  $\vec{y}_1, \dots, \vec{y}_{|R|} \sim D$   
In this talk,  $D = \text{Uniform } ([0,1]^d)$

Connect  $i \in L, j \in R$  if  $\parallel \vec{x}_i - \vec{y}_j \parallel_{\infty} \leq \Delta$

## Applications:

- Dynamic spatial matching  
[Kanoria '23]
- Opinion dynamics  
[Zhang et al. '14]
- Contagion in wireless networks  
[Nekovee '07]

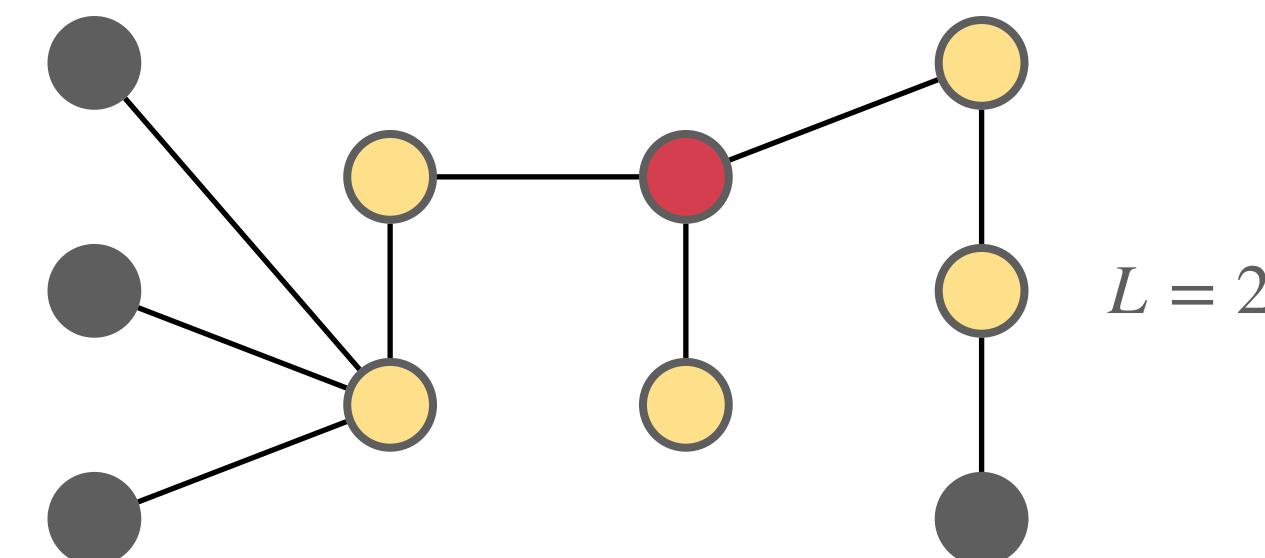
# GNNs learn local graph functions

E.g., Tahmasebi et al., AISTATS'23

A function  $f$  over graphs is  **$r$ -local** if

$f(G)$  only relies on info in the  **$r$ -hop neighborhood**  $N_r(v)$  of each node  $v$   
i.e., exist functions  $\phi, \psi$  such that  $f(G) = \phi \left( \psi \left( N_r(v) \right)_{v \in V} \right)$

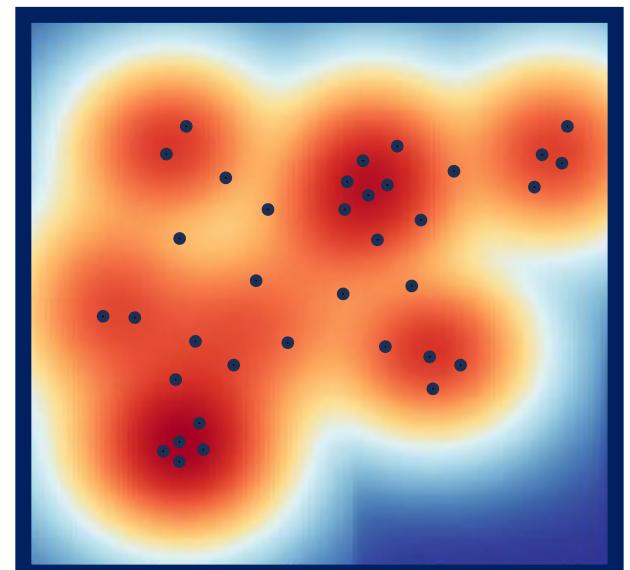
**Observation:** GNNs with  $L$  layers compute  $L$ -local functions



# Value-to-go is locally approximable

Function  $f$  is  **$(r, \epsilon, \delta)$ -locally approximable** over a graph distribution  $\mathcal{G}$  if:  
there exists an  $r$ -local function  $h$  such that  $|f(G) - h(G)| \leq \epsilon f(G)$   
with probability  $1 - \delta$  over  $G \sim \mathcal{G}$

**Main theorem (informal):** For  $\epsilon, \delta \in (0, 1)$  and  $\Delta = O\left((|L| + |R|)^{-1/d}\right)$ ,  
VTG is  $(\log(|L| + |R|), \epsilon, \delta)$ -locally approximable over  $\mathcal{G}(|L|, |R|, D, \Delta)$



*Proof intuition:*

- Graphs can be split into “clusters” with few inter-cluster edges
- Matching within clusters is nearly optimal

# Outline

1. Introduction
2. Graph neural networks (GNNs) for online matching [ICML'24]
  - i. Background
  - ii. GNN setup
  - iii. Theoretical guarantees
  - iv. Experiments**

# Graph datasets

## Random graph families

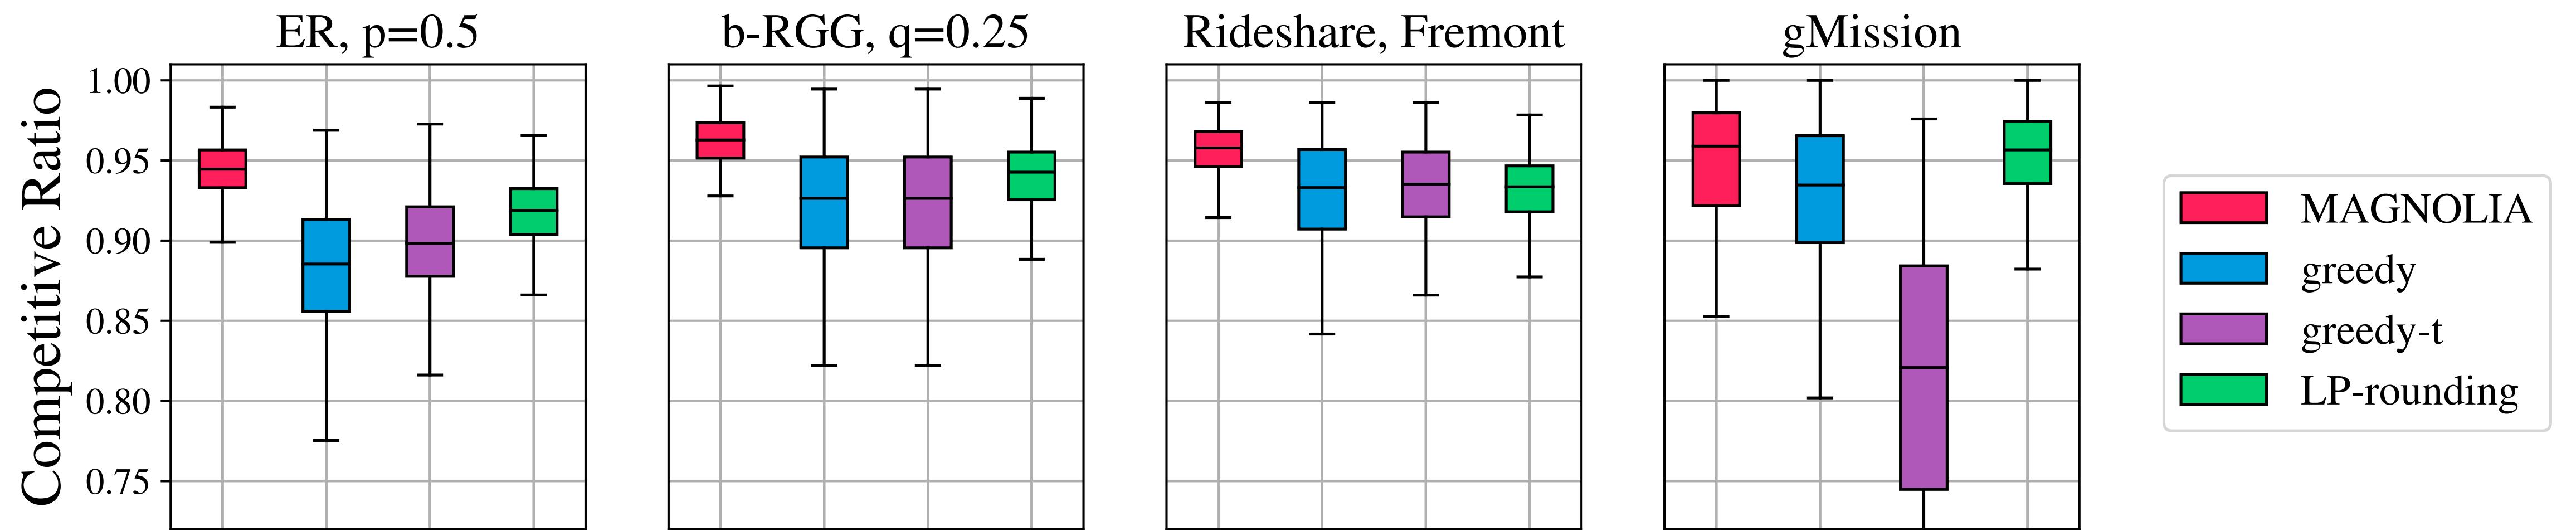
- Erdős-Rényi graphs with  $U(0,1)$  edge weights
- Barabási-Albert graphs with  $U(0,1)$  edge weights
- Random geometric graphs



## Semi-synthetic graphs

- gMission: crowdsourcing data for assigning workers to tasks  
[Chen et al., VLDB'14]
- OSMnx: Python library for analyzing street networks  
[Boeing, Comput. Environ. Urban Syst. '17]

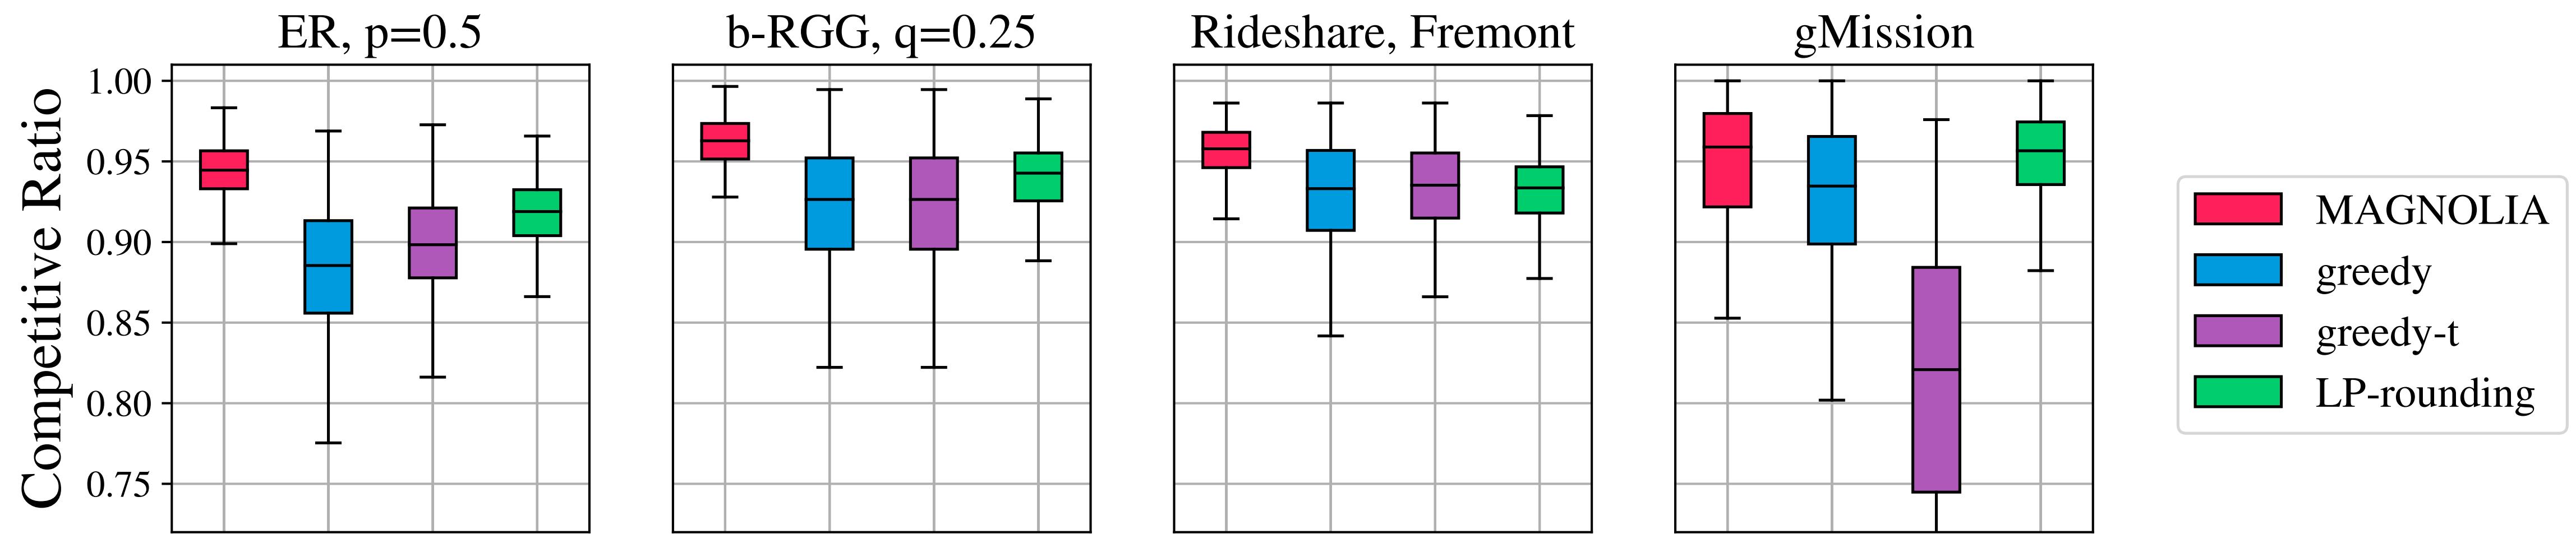
# Experiments



Competitive ratio:  $\frac{\text{ALG}}{\text{Offline optimal}}$

Trained on graphs of size (6x10), competitive ratios shown for graphs of size (10x30)

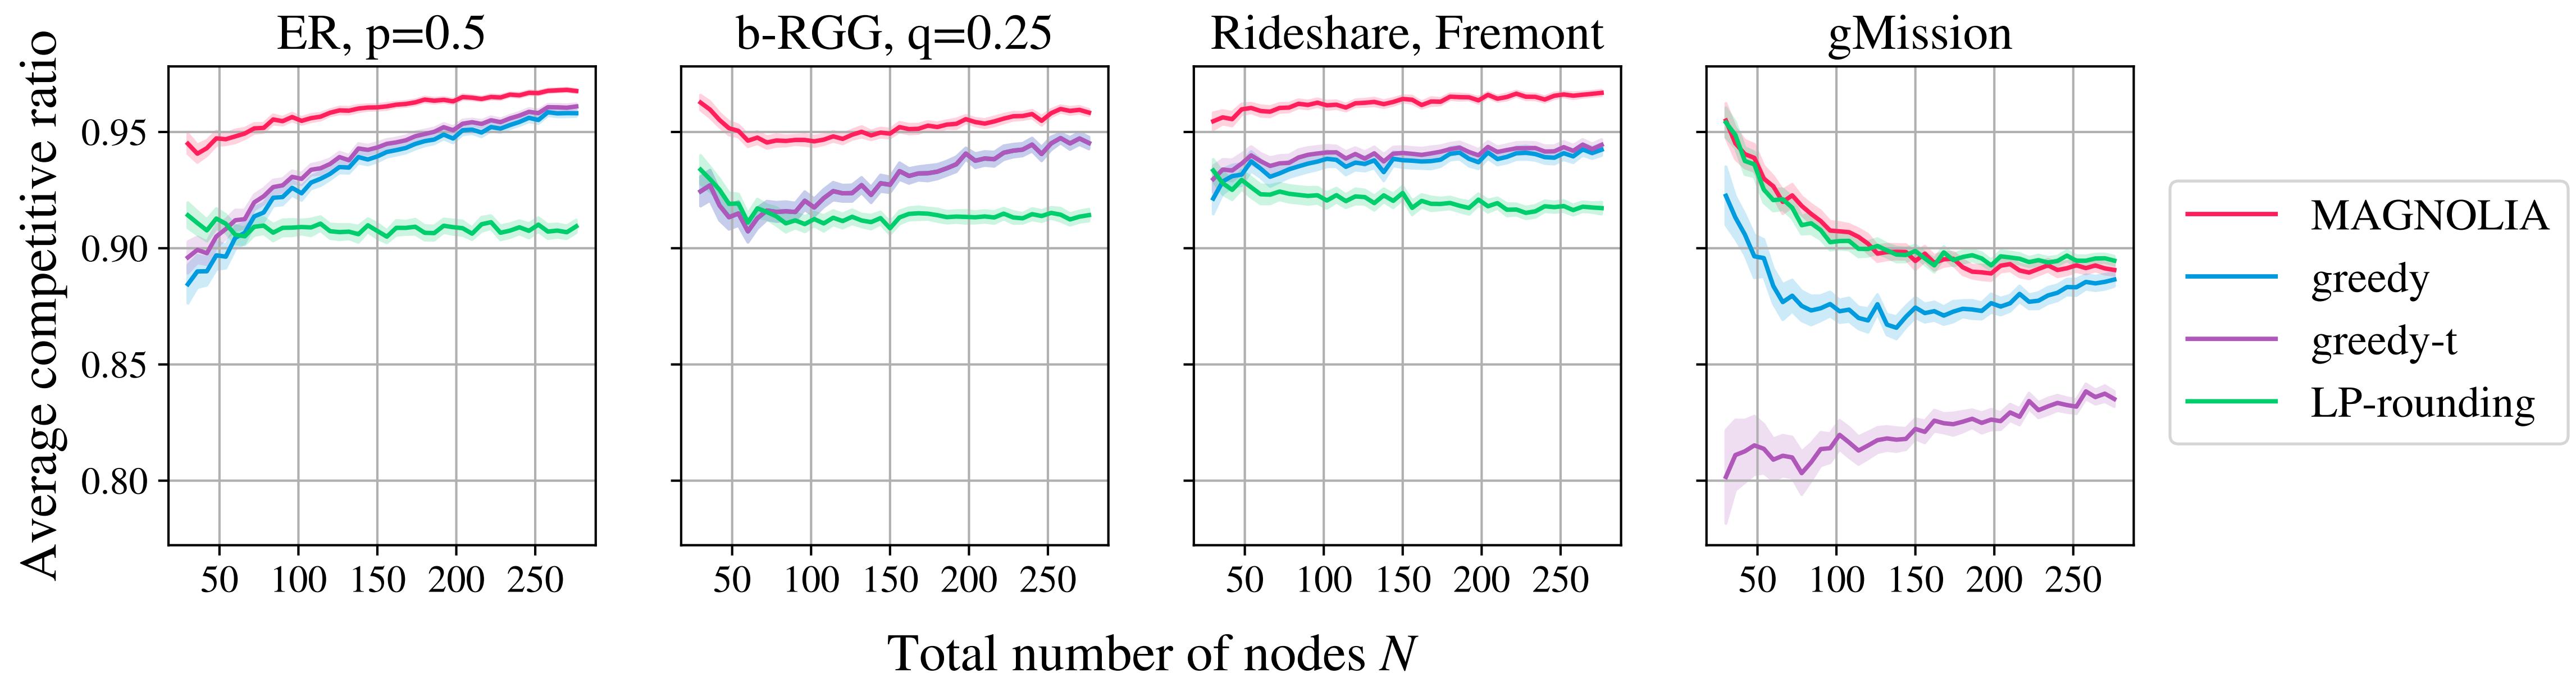
# Experiments



- greedy: selects highest weight available edge
- greedy-t: selects greedy edge if weight exceeds threshold, otherwise skips
- LP-rounding: pen-and-paper 0.632-approximation to optimal online algorithm

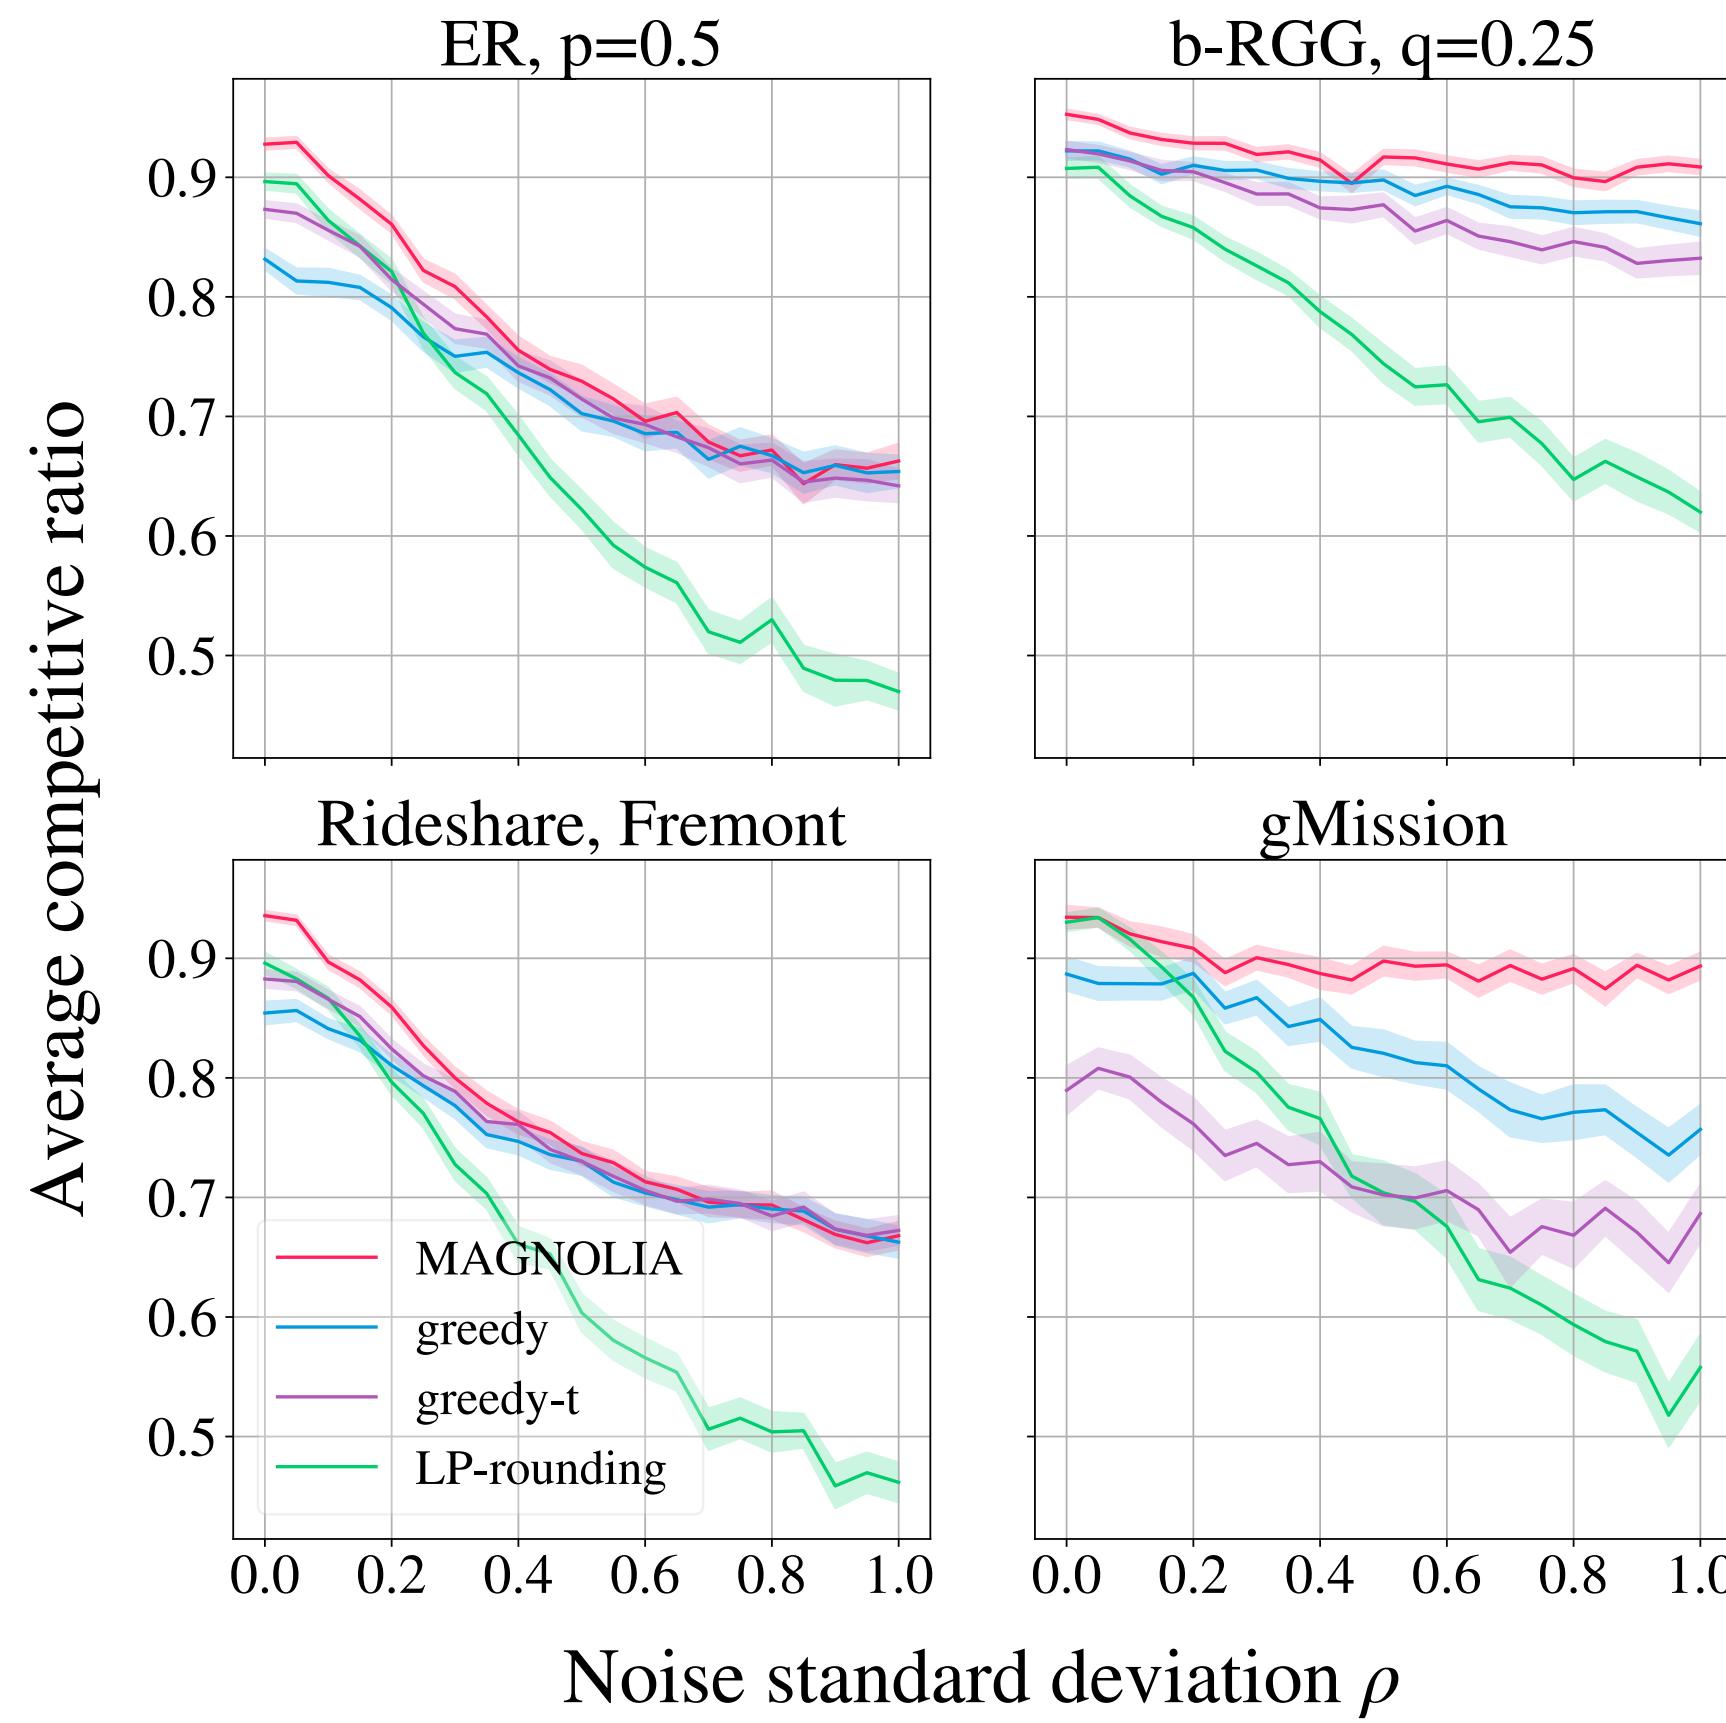
Papadimitriou et al. [EC'21], Braverman et al. [EC'22], Naor et al. ['23]

# Experiments: Size generalization



Performance generalizes to graphs up to 20x larger than training graphs (6x10)

# Experiments: Noise



Weights, arrival probabilities may be unknown  
Instead, only **noisy estimates** from data

Add  $N(0, \rho^2)$  noise independently to  
edge weights and arrival probabilities

MAGNOLIA **more robust** than baselines

# Conclusions and future directions

## Graph neural networks (GNNs) for online matching

- NP-hard problem, applications in online advertising, crowdsourcing, ...
- Supervise using (exponential time) online optimal policy
  - Train on small graphs, test on large graphs
- Theoretical guarantees to understand why GNNs are suitable for this task

## Future directions

- **More unknowns:** graph structure, edge weights, arrival probabilities, ...
- **Other problems:** locality analysis might apply to other problems

# **Online matching with graph neural networks**

**Ellen Vitercik** (Stanford)